

Exploring the hierarchical structure of human plans via program generation

Carlos G. Correa, Sophia Sanborn, Mark K. Ho, Frederick Callaway, Nathaniel D. Daw, Thomas L. Griffiths

Presenter: Hanqi Zhou, 2024.03.19

Goal

- Exploring the hierarchical structure of human plans via program generation.

Goal


- Exploring the hierarchical structure of human plans via program generation.
- Task: make tomato sauce for dinner
- Plan:
 - prepare the tomatoes →
 - place the tomatoes into a large oven or saucepan



Goal

- Exploring the hierarchical structure of human plans via program generation.
- Task: make tomato sauce for dinner

- Plan:



place the tomatoes into a large oven or saucepan

Goal

- Exploring the hierarchical structure of human plans via program generation.
- Task: make tomato sauce for dinner
- Plan:

```
def make_tomato_sauce (N):  
    i = 0  
  
    ingredient = None  
  
    while i < N:  
        ingredient += rep ( prep_tomato, N )  
        i ++  
  
    place_oven (ingredient)
```

```
def prep_tomato ():
```

```
    ...
```

```
def place_oven (x):
```

```
    ...
```

```
def rep (x, y):
```

```
    ...
```

Background - Bayesian Program Induction

- Infer programs (*actions*) π consistent with some observed data (*task*) d

$$p(\pi \mid d) \propto p(d \mid \pi)p(\pi)$$

- Why **Bayesian induction** and why **program**?
 - Compositionally, causality and learning to learn (Lake et al., 2015)

Assumptions

- Assumption:
 - Human behavior is inherently **hierarchical**, resulting from the decomposition of a task into subtasks or an abstract action into concrete actions.
- Method
 - Utilized a process-tracing experiment with the **Lightbot** game, applying **Bayesian program induction** to model and analyze hierarchical plan formation.

Problem setting

The interface displays a robot on a 5x5 grid of grey blocks. Three blue squares are located at (1,2), (2,3), and (3,4) in a 0-indexed coordinate system where (0,0) is the top-left. The robot is currently at (1,0). Above the grid are buttons for 'Run', 'Quick Run' (with a lightning bolt icon), and a 'Clear Main' button. To the right of the grid is an 'Instruction Count' box showing the number 13. Further right is a table for process tracing:

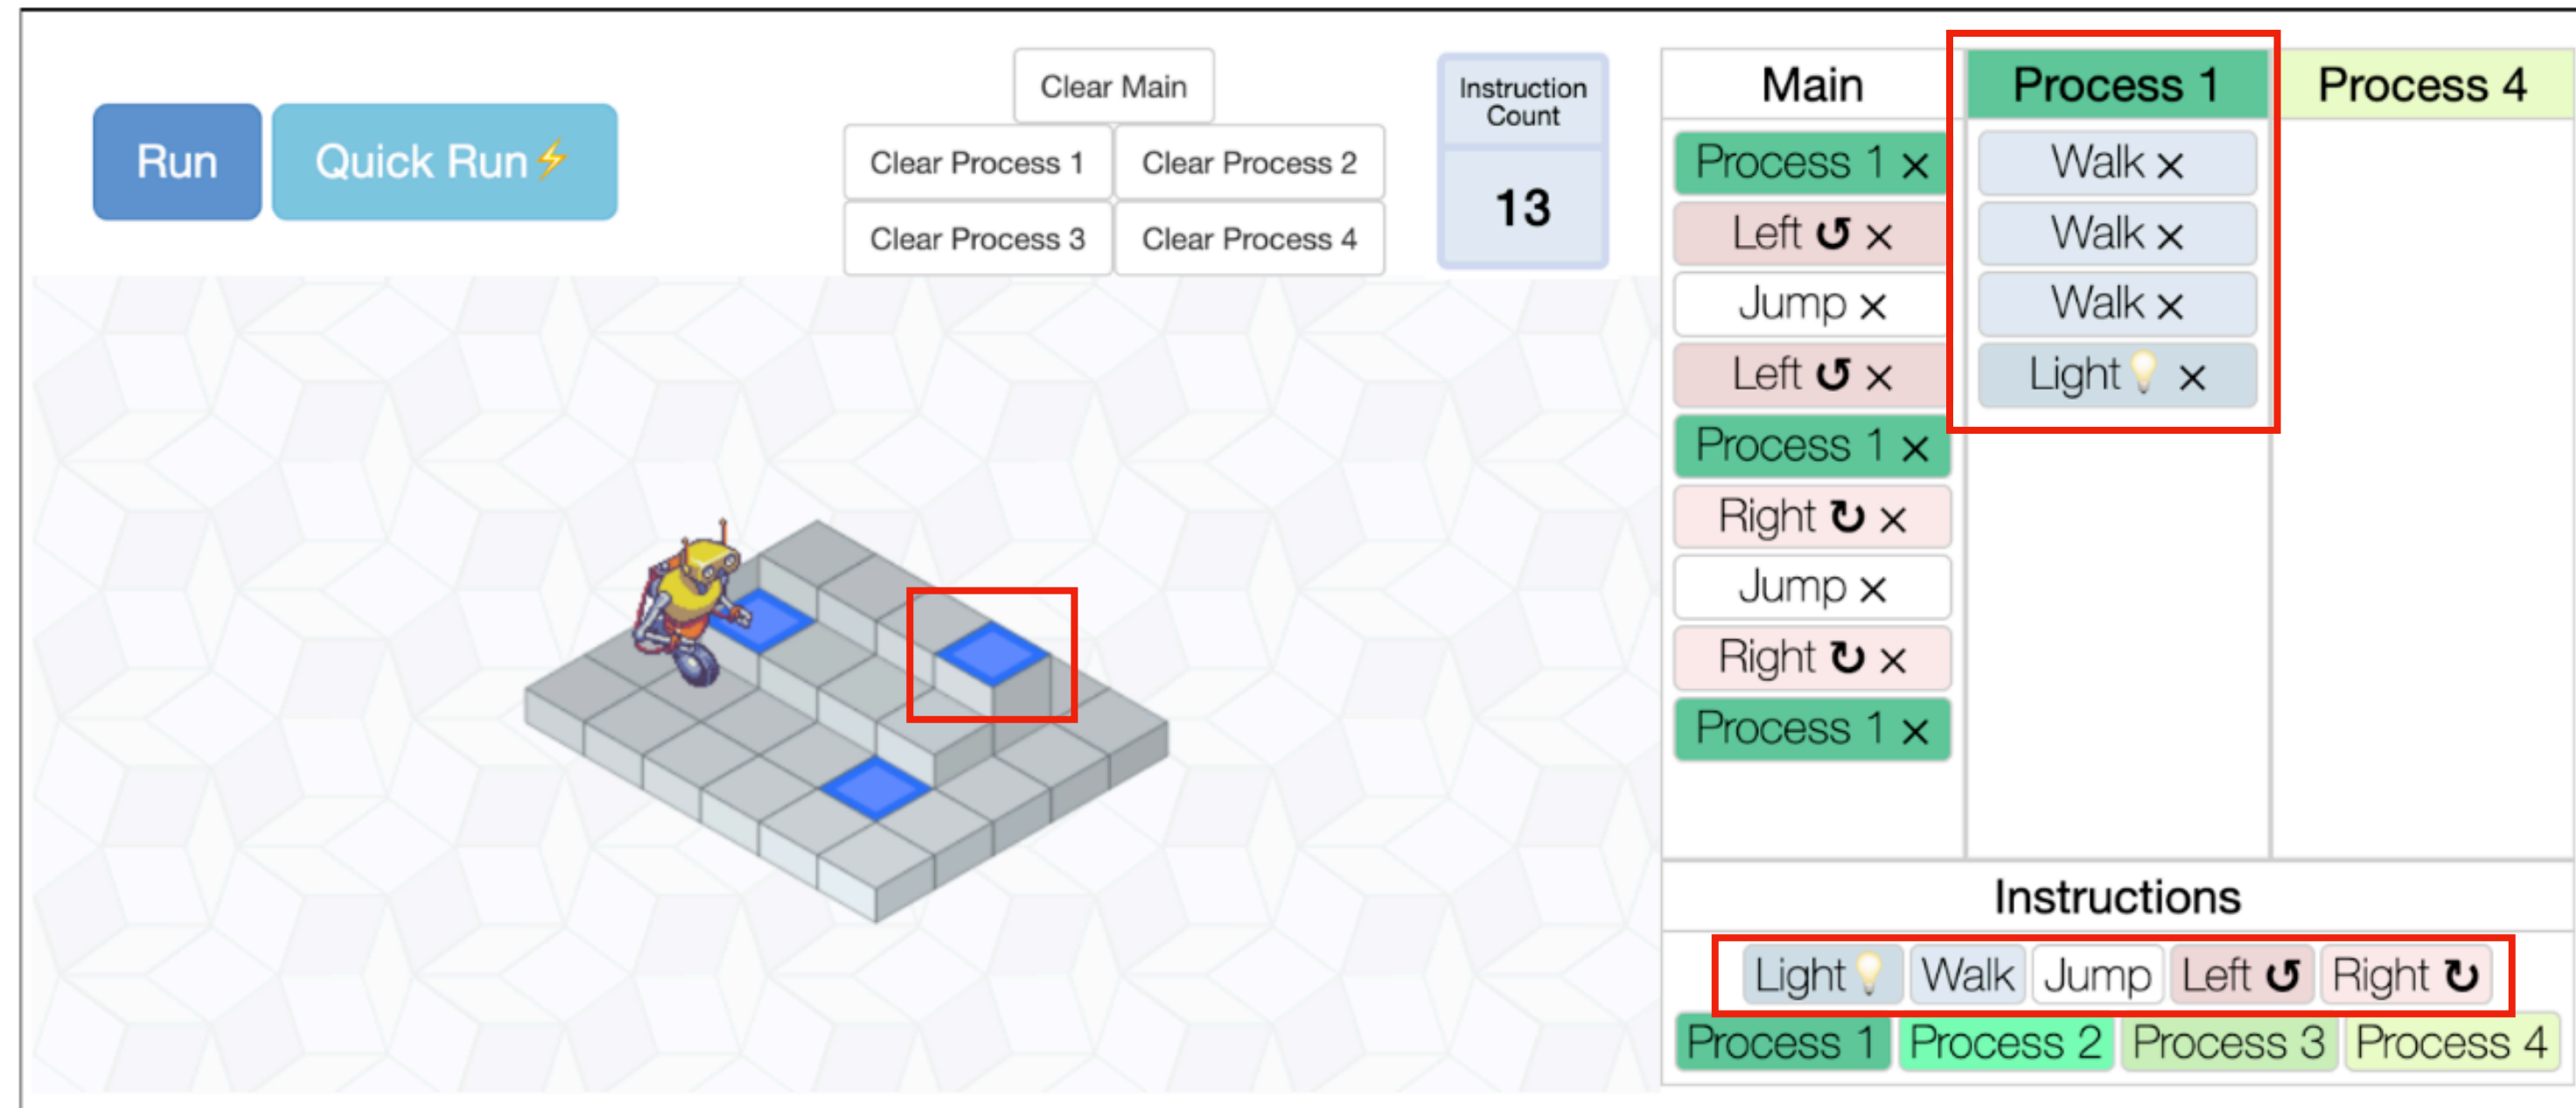
Main	Process 1	Process 4
Process 1 x	Walk x	
Left ↶ x	Walk x	
Jump x	Walk x	
Left ↶ x	Light 💡 x	
Process 1 x		
Right ↷ x		
Jump x		
Right ↷ x		
Process 1 x		

Below the table is an 'Instructions' palette with buttons for 'Light 💡', 'Walk', 'Jump', 'Left ↶', and 'Right ↷'. At the bottom of the interface are four colored tabs: 'Process 1' (green), 'Process 2' (teal), 'Process 3' (light green), and 'Process 4' (yellow).

Lightbot - process tracing

- Goal: activate all blue squares
- Action set
 - Primitives
 - Subroutines/processes
- Execution trace
 - A deterministic sequence of actions

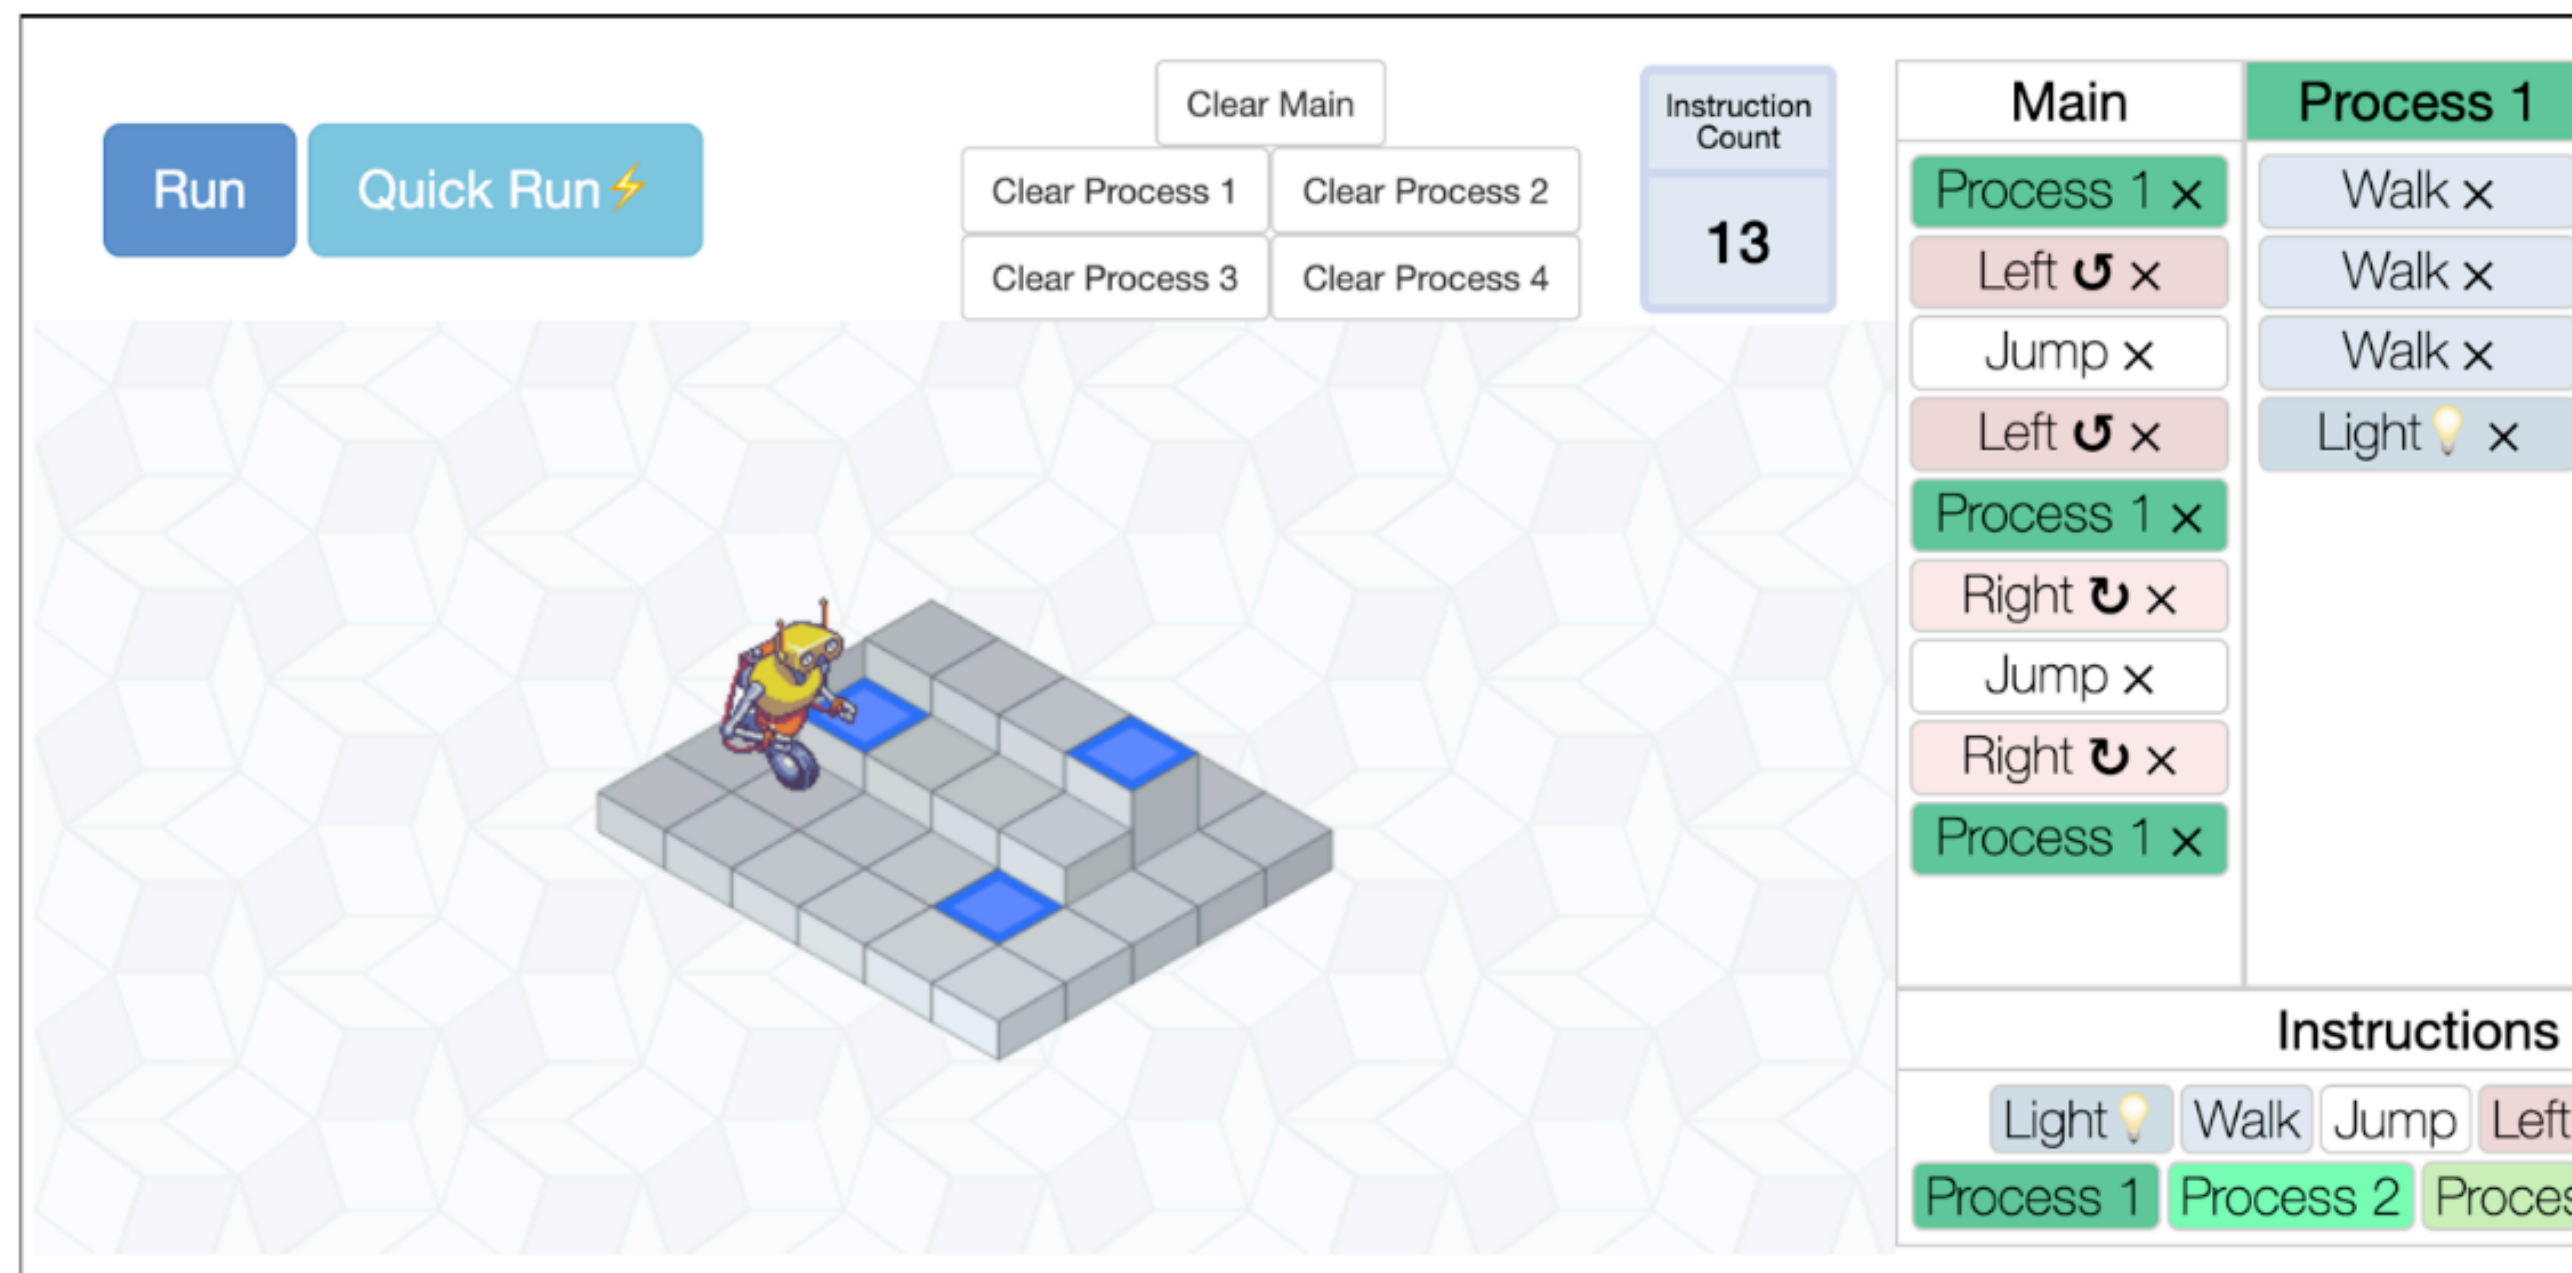
Problem setting



Lightbot - process tracing

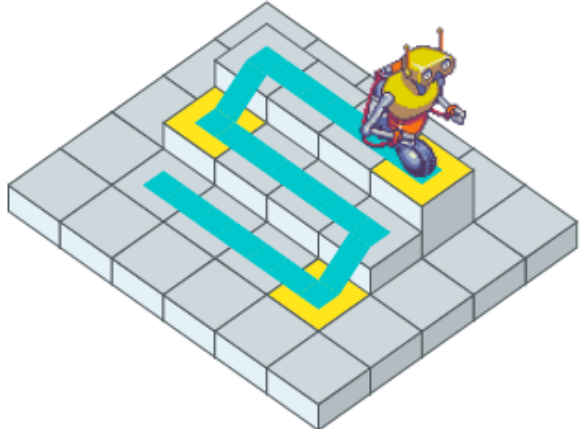
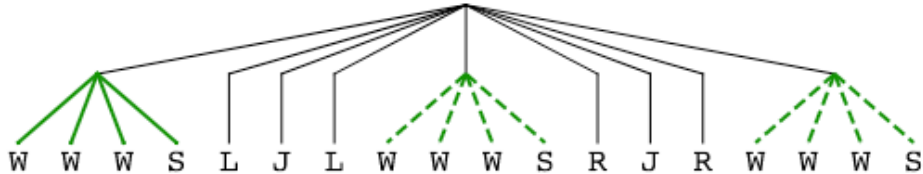
- Goal: activate all blue squares
- Action set
 - Primitives
 - Subroutines/processes
- Execution trace
 - A deterministic sequence of actions

Model of the environment - MDP



- An undiscounted, deterministic Markov Decision Process (MDP)
 - State set \mathcal{S} : environment state (active or not), agent state (location & orientation)
 - Goal states $\mathcal{G} \subset \mathcal{S}$
 - Action set \mathcal{A} : five primitive actions
 - Transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$
 - Reward function $R : \mathcal{S} \times \mathcal{S} \rightarrow \mathbb{R}$:
 $R(s, g) = 1$ only if $g \in \mathcal{G}$

Model of the agent - Program

Trace	Program	Step Count	Program Length
		18	13

- Program $\pi = \{\rho^0, \rho^1, \dots\}$
- Subroutine $\rho^i = \{\rho_0^i, \rho_1^i, \dots\}$
- Instructions $\rho_j^i \in \mathcal{A} \cup \{\rho^1, \rho^2, \dots\}$

Run

Quick Run ⚡

Clear Main

Clear Process 1

Clear Process 2

Clear Process 3

Clear Process 4

Instruction Count

13

Main

Process 1 x

Left ⤴ x

Jump x

Left ⤴ x

Process 1 x

Right ⤵ x

Jump x

Right ⤵ x

Process 1 x

Process 1

Walk x

Walk x

Walk x

Light 💡 x

Process 4

Instructions

Light 💡 Walk Jump Left ⤴ Right ⤵

Process 1 Process 2 Process 3 Process 4


Model - Bayesian Program Inference

$$p(\pi \mid \text{success}) \propto p(\text{success} \mid \pi) p(\pi)$$

- Baseline models
 - Prior:
 - Random $\log p(\pi) = \frac{1}{|\Pi|}$
 - Step count $\log p(\pi) \propto -|\tau(\pi)|$
 - Program/Description length $\log p(\pi) \propto -\text{DL}(\pi)$
 - Reuse-priority
 - Likelihood: binary

Model - Program Prior towards Reuse

- Key idea: past use of subroutines should inform the probability assigned to future use of subroutines.
- Method: Chinese Restaurant Process (CRP)



$$p(\rho^i) = p_{\text{end}} p(\rho_0^i = a) \prod_{j=1} (1 - p_{\text{end}}) p(\rho_j^i = a)$$

$$p(\rho^{n+1} \mid \rho^{1:n}) = \begin{cases} \frac{\alpha}{n + \alpha} & \text{if } \rho^{n+1} = m + 1 \\ \frac{n_k}{n + \alpha} & \text{if } 1 \leq \rho^{n+1} = k \leq m \end{cases}$$

Experiments

- Participants
 - have to solve 10 tasks
 - are incentivized to write short programs (to encourage subroutine use)

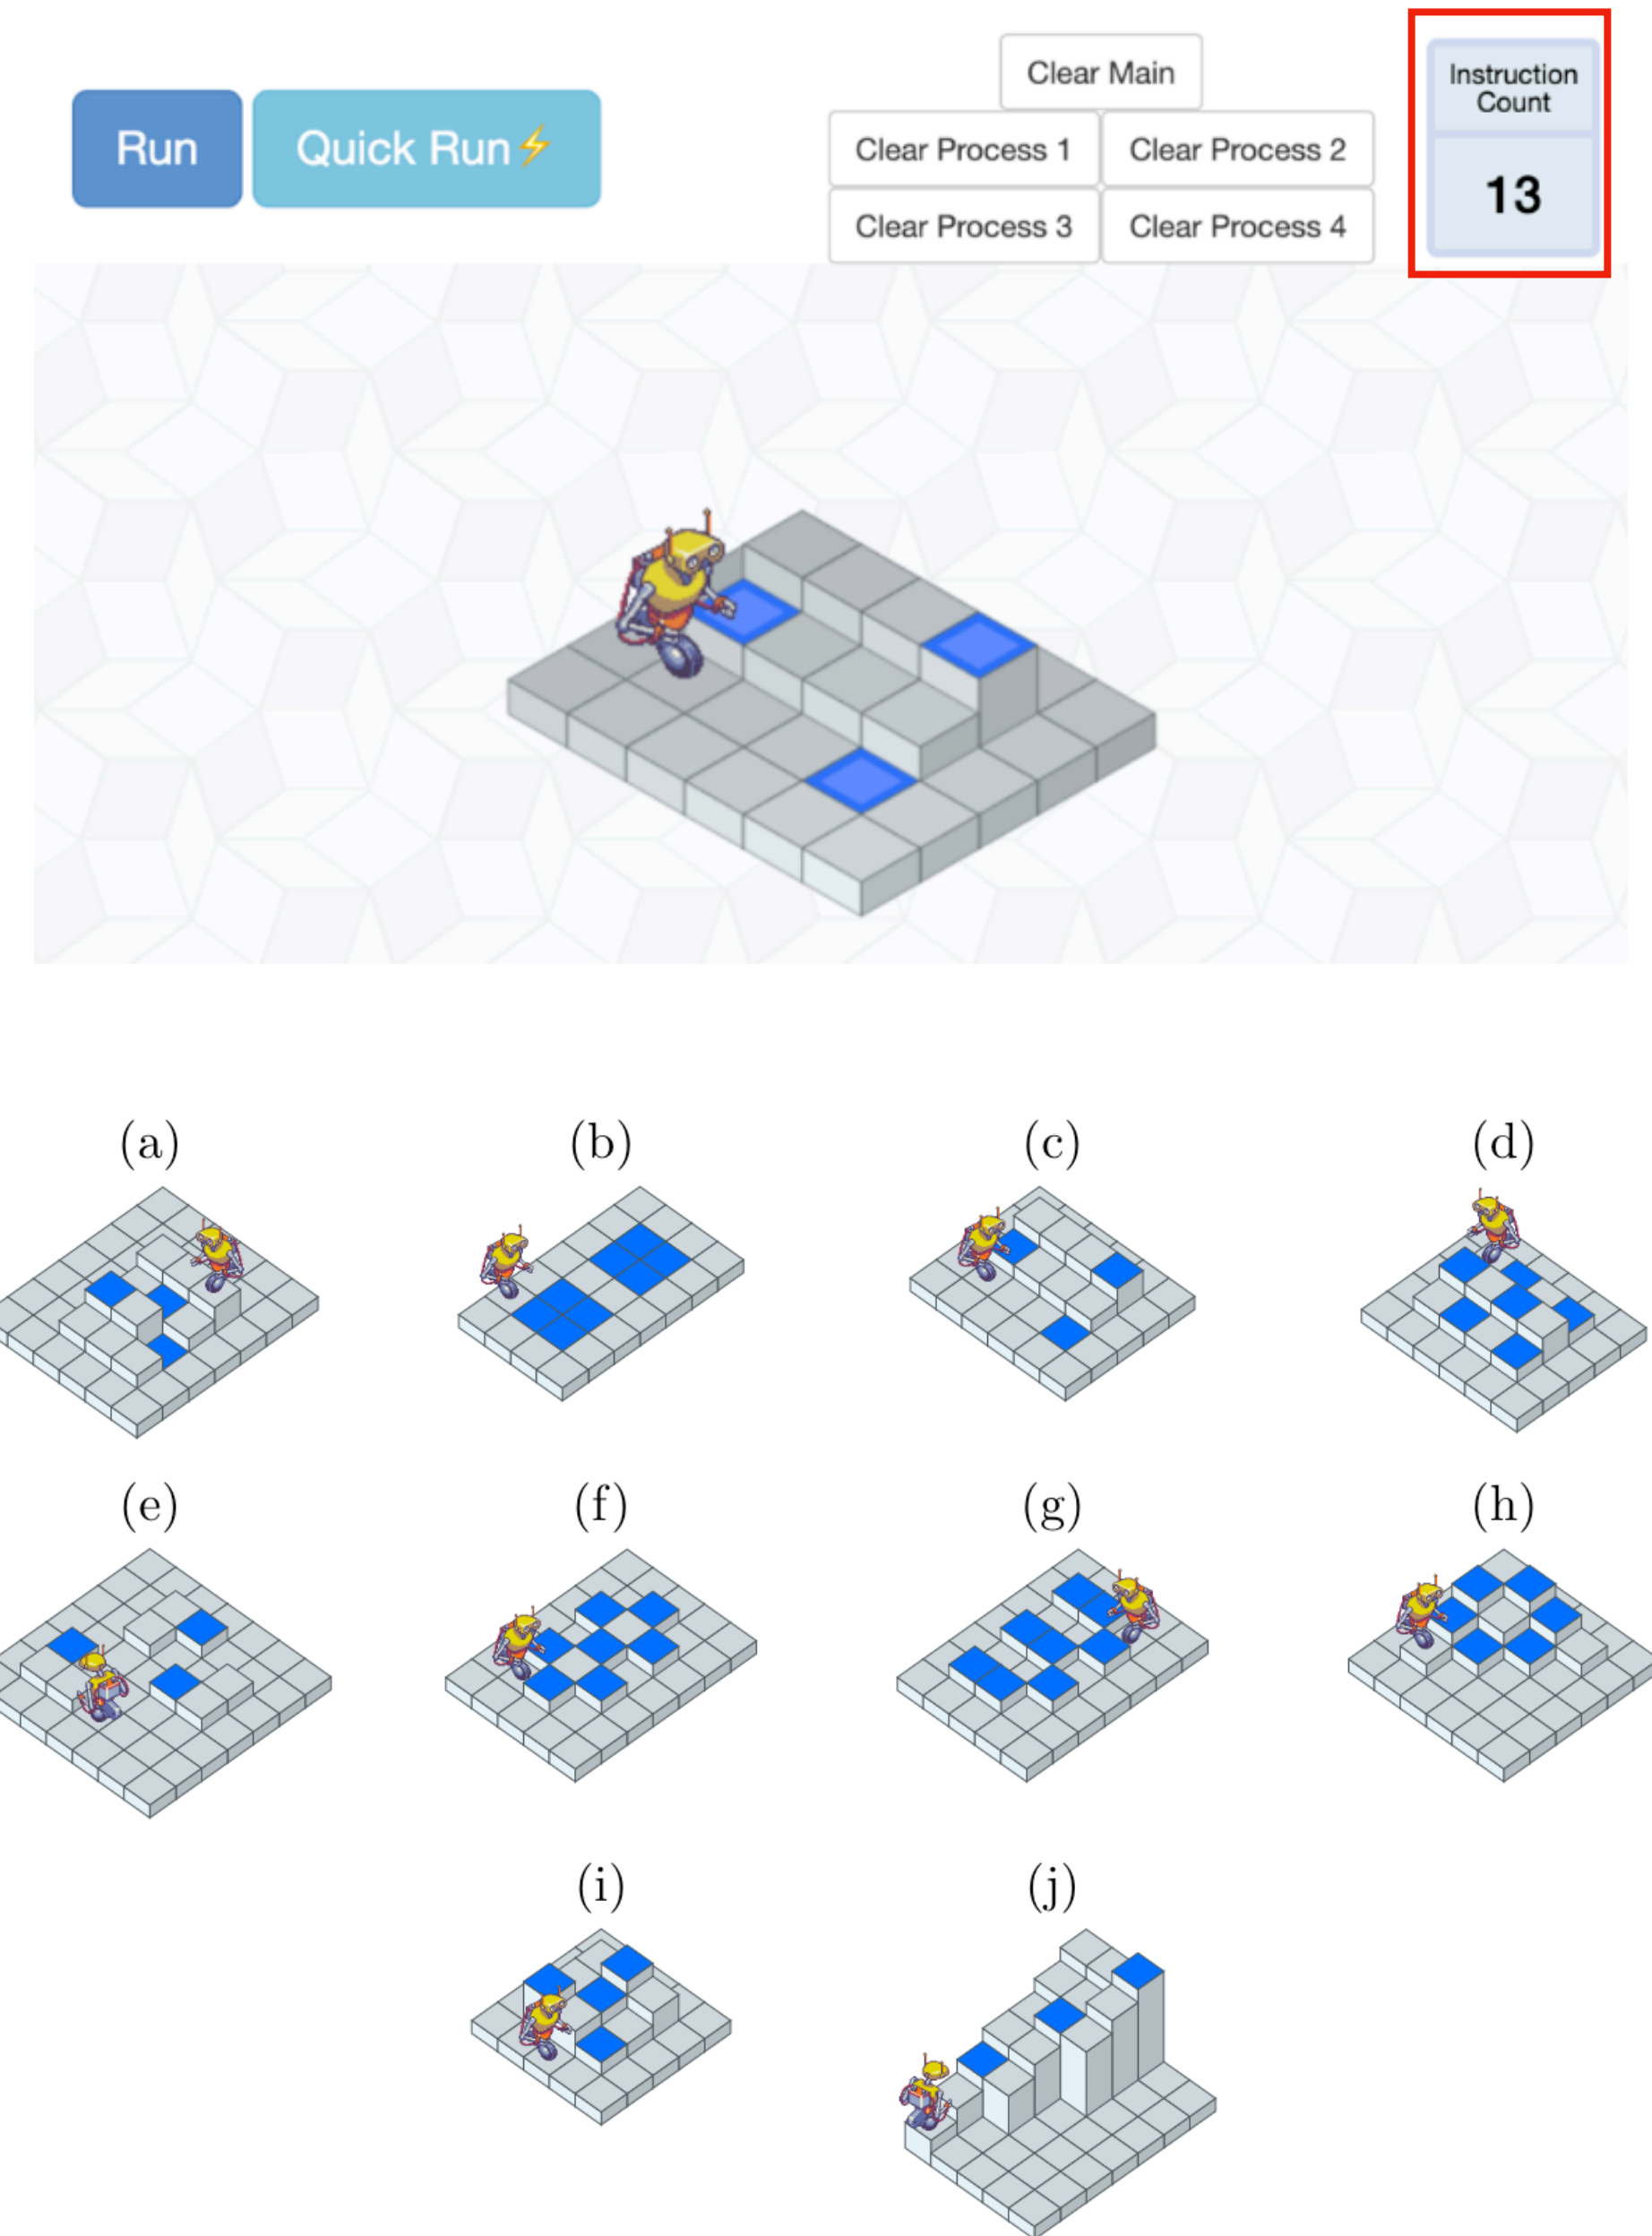




Figure 3: The tasks participants completed in the experiment.

Results

The step cost model is a better fit to behavior than the MDL



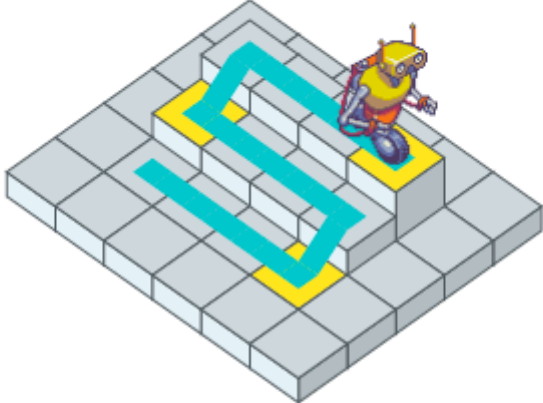
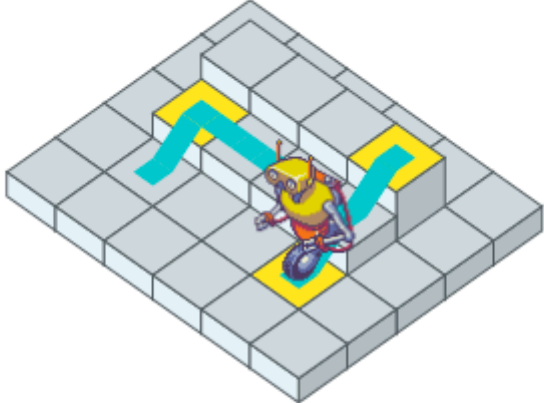
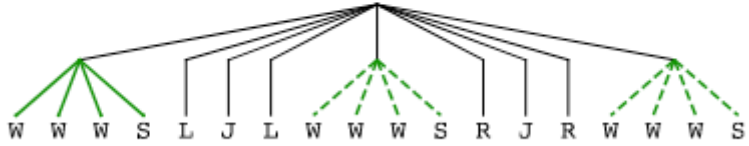
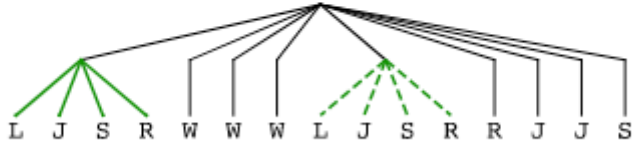
(a)



(b)

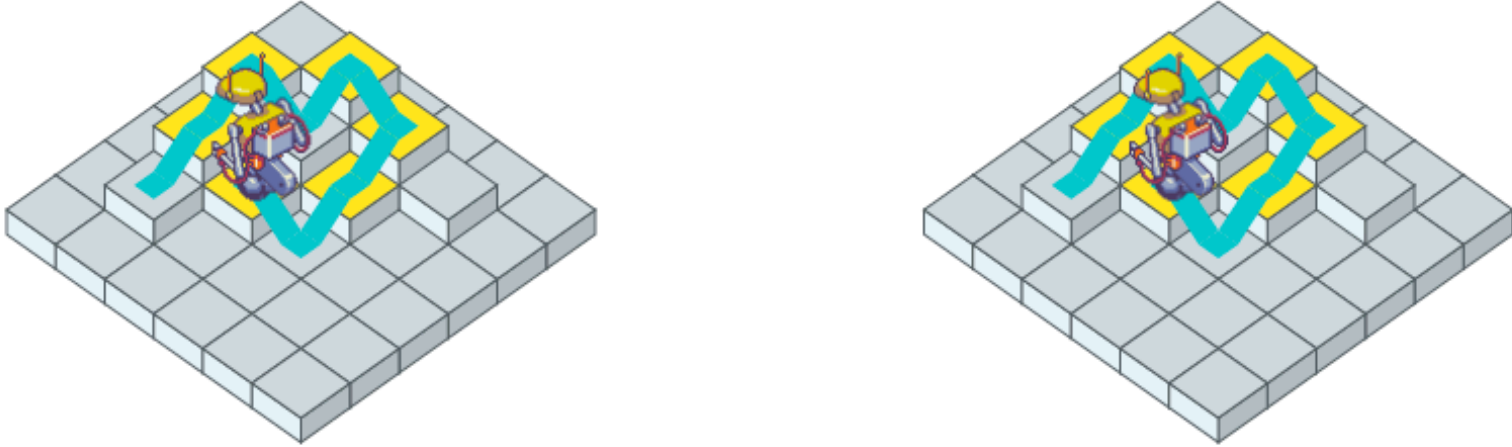
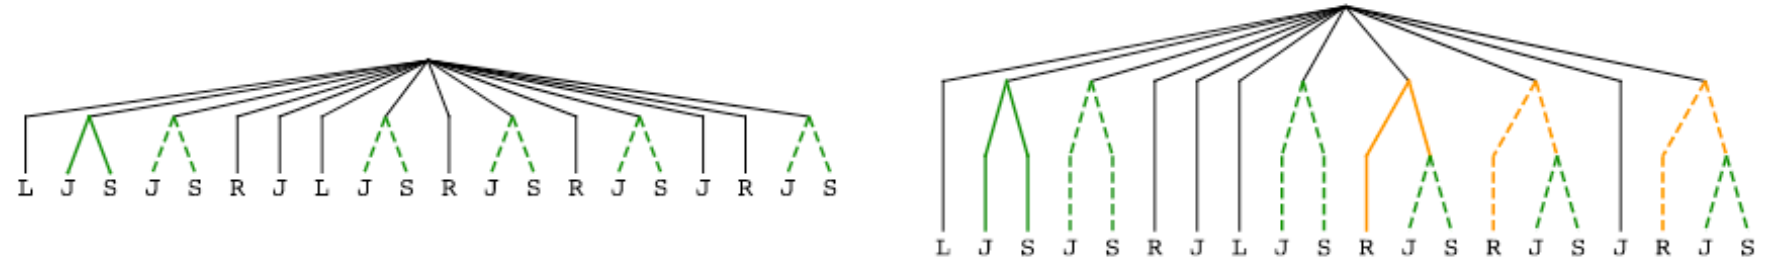
Results

Participant programs are inconsistent with alternative accounts

trace		
program		
participant count	59	1
step count	18	15
program length	13	13
grammar induction prior (log)	-31.97	-33.17

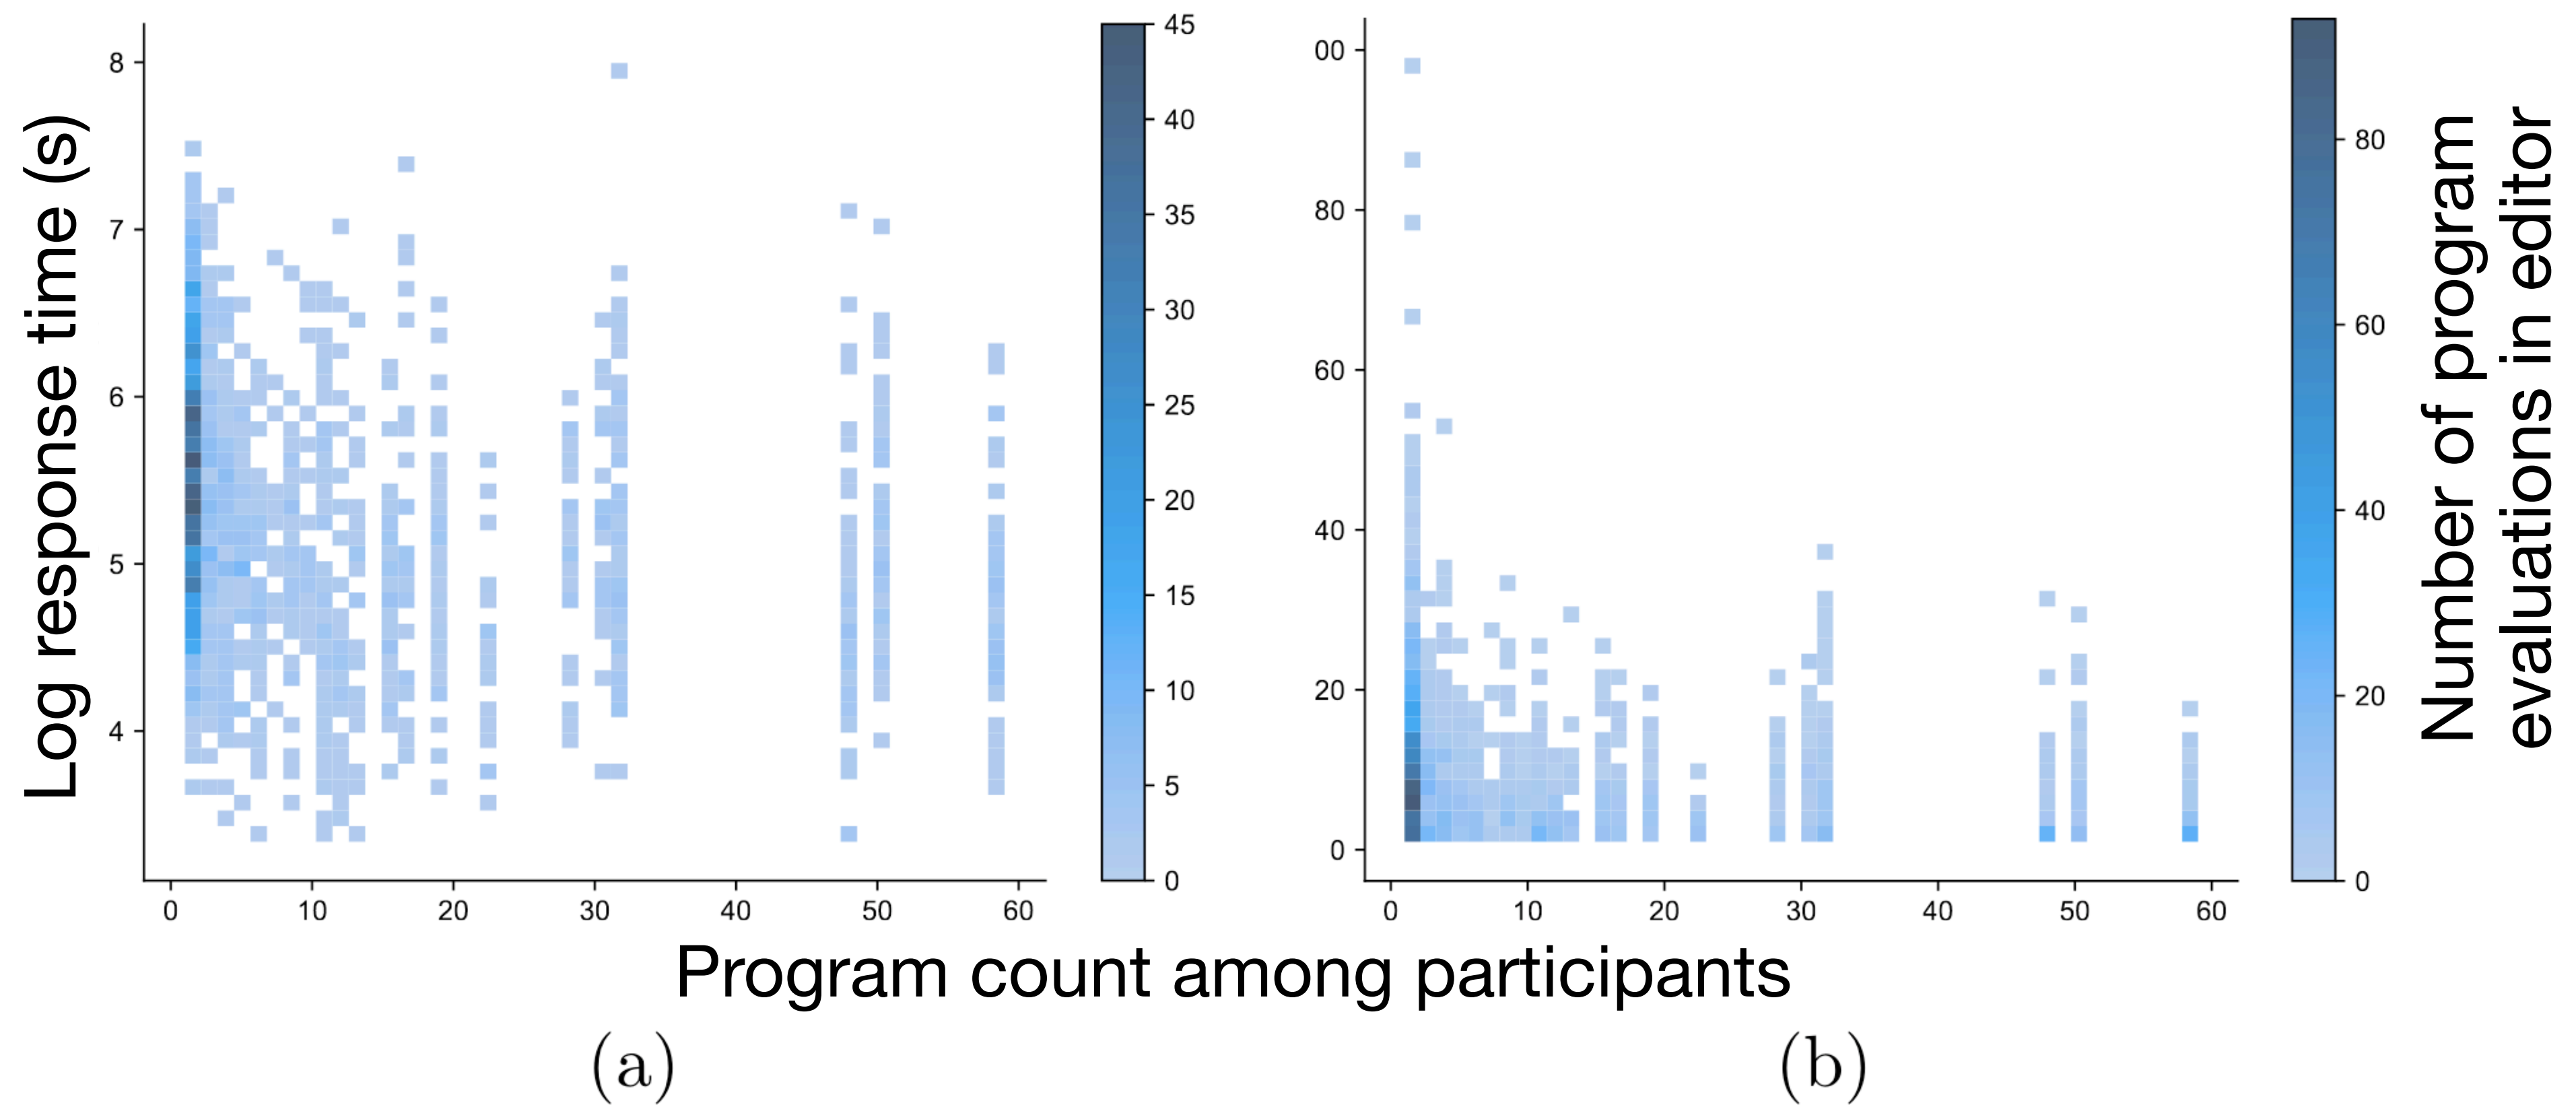
Results

Participant programs are inconsistent with alternative accounts

trace		
program		
participant count	19	3
step count	20	20
program length	16	15
grammar induction prior (log)	-35.06	-37.49

Results

Common programs are easier to write



Findings

- In this work, we find evidence that people favor both shorter programs and those with fewer actions.
- However, we additionally identify examples in which people prefer reusing subroutines above and beyond the predictions of either account.
- We formalize this preference as a prior belief that subroutine use should be biased towards subroutines that were used often in the past.